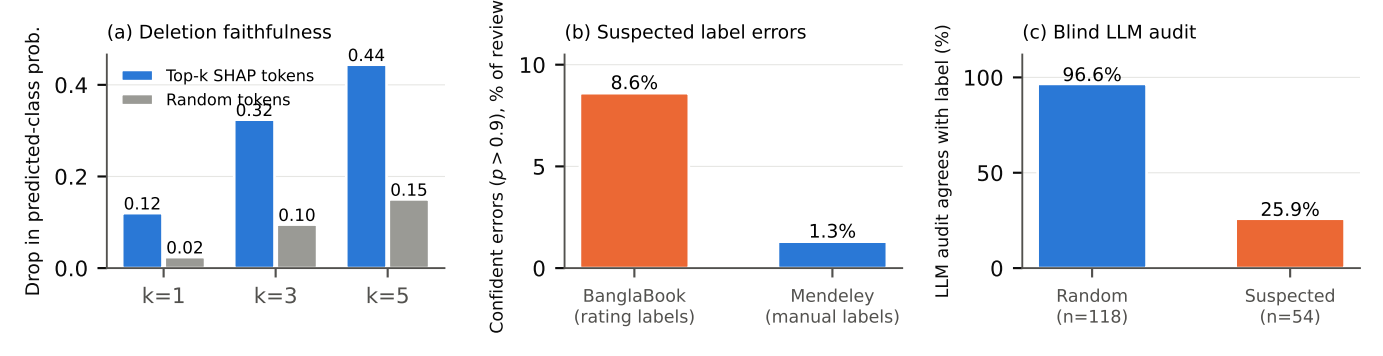
(a) Deletion faithfulness
Drop in predicted-class prob.
Top-k SHAP tokens
Random tokens
0.12
0.02
0.32
0.10
0.44
0.15
0.0
0.2
0.4
k=1
k=3
k=5
(b) Suspected label errors
Confident errors ($p > 0.9$), % of reviews
8.6%
1.3%
0
5
10
BanglaBook
(rating labels)
Mendeley
(manual labels)
(c) Blind LLM audit
LLM audit agrees with label (%)
96.6%
25.9%
0
50
100
Random
(n=118)
Suspected
(n=54)